

# Machine Learning

## AIAA 5046, Spring 2024

Sihong Xie

AI Thrust

HKUST-GZ

Goals:

- MDP
- Value functions

Readings:

RL Chap 3-6

# Sequential decision making

Sequential decision making:

- an agent interacts with an environment, and
- takes a sequence of actions to reach a goal.

Different from supervised learning:

- data are generated via interactions and can change with the agent's behavior;
- data are not I.I.D.;
- no one-shot prediction and immediate feedback.

Different from unsupervised learning:

- finding clustering patterns can't solve sequential decision making problems.
- but can help find representation of the environment to help.

Drone control



Game




# Markov Decision Processes

Markov Decision Processes: a mathematical description of sequential decision making.

- We use  $t = 1, 2, \dots$  to denote the steps of decision-making.

An MDP has five components.


- A set of states  $\mathcal{S}$  and current state  $S_t \in \mathcal{S}$
- Action space  $\mathcal{A}(S_t)$
- Reward  $R_t$
- Dynamics  $\Pr(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$
- Discounting factor  $0 \leq \gamma \leq 1$



# Markov Decision Processes

Learning from interactions:

The agent interacts with the environment and learns from the interaction experiences.



A trajectory or experience is denoted as

$$S_0, A_0, R_0, S_1, A_1, R_1, S_2, A_2, R_2, \dots$$

These are all random variables, as the agent selects actions stochastically, and the MDP dynamics (rewards and state transition) are stochastic.


# Markov Decision Processes

Goal of reinforcement learning:

- maximizing the expectation of the discounted cumulative rewards (called ``return''):

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

( $\gamma$  can ensure convergence of the series)




An example trajectory:

State, Action, reward  
(1,1), Up, -0.04,  
(1,2), Up, -0.04,  
(1,3), Right, -0.04,  
(2,3), Right, -0.04,  
(3,3), Right, -0.04,  
(4,3), NoAction, 1.

Return = 0.8 with  $\gamma=1$ .

The following MDP shows that it is important to maximize return, not immediate reward. Starting from state u, going left has a higher immediate reward than going right, which can lead to a high reward when reaching z.



# Policy



The agent interacts with the environment using a policy, which decides how to choose an action in a state.

$$0 \leq \pi(a|S_t) \leq 1, \sum_{a \in \mathcal{A}(S_t)} \pi(a|S_t) = 1$$

- Example on the right:
  - a state is a position in the maze,
  - an action is a direction to go next.

A deterministic sequence of actions will fail in the face of environment uncertainty:

- there is a non-zero chance that the fixed action sequence (U, U, R, R, R) can lead to the undesirable state (4,2).



# State-value functions


To optimize the policy, a policy is first evaluated by the state-value functions:

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$$

More explicitly,

$$\begin{aligned} v_\pi(s) &= \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] \\ &= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s'} \sum_r \Pr(s', r | s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} | S_{t+1} = s']] \\ &= \sum_a \pi(a|s) \sum_{s'} \sum_r \Pr(s', r | s, a) [r + \gamma v_\pi(s')] \\ &= \sum_a \sum_{s'} \sum_r \pi(a|s) \Pr(s', r | s, a) [r + \gamma v_\pi(s')]. \end{aligned}$$

Expectation taken over all possible trajectories sampled according to the policy and the environment dynamics.




Bellman equation:  $v_\pi(s) = \sum_r \Pr_\pi(r|s)r + \gamma \sum_{s'} \Pr_\pi(s'|s)v_\pi(s')$

# State-value functions

For one policy  $\pi$ , there is one state-value function  $v_\pi(s)$ , defined on the state space  $\mathcal{S}$ .

A special case:  
each cell has  
an action  
according to a  
*deterministic* policy.



3	0.812	0.868	0.918	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388


each cell  $s$   
has  $v_\pi(s)$

# Action-value function

A related value function: expected return when taking action  $a$  at state  $s$ .

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a] = \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s, A_t = a]$$

$$\begin{aligned} q_\pi(s, a) &= \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s, A_t = a] \\ &= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a] \\ &= \sum_{s'} \sum_r \Pr(s', r | s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} | S_{t+1} = s', S_t = s, A_t = a]] \\ &= \sum_{s'} \sum_r \Pr(s', r | s, a) [r + \gamma \sum_{a'} \pi(a' | s') q_\pi(s', a')] \\ &= \sum_r \Pr(r | s, a) r + \gamma \sum_{s'} \sum_{a'} \Pr(s' | s, a) \pi(a' | s') q_\pi(s', a'). \end{aligned}$$



backup diagram

# Optimal policies


The agent wants to find the optimal policy  $\pi^*$ , defined as

$$v_{\pi^*}(s) \geq v_{\pi}(s), \quad \forall s \in \mathcal{S}, \quad \forall \pi.$$

Bellman optimality equation

$$v_{\pi^*}(s) = \max_{a \in \mathcal{A}(s)} \sum_{s', r} \Pr(s', r | s, a) [r + \gamma v_{\pi^*}(s')]$$

Intuitively, no matter where the agent is, the optimal policy always leads to the best expected return.



At any state  $S$ , the optimal policy must select the best action, then follow the optimal policy from the successor state  $S'$ , considered as a subproblem. The value function  $v_{\pi^*}(s')$  caches an optimal values of the subproblems. The optimal action selection stores an optimal solutions.


# Optimal value function

Bellman optimality equation for the action-value function

$$q_{\pi^*}(s, a) = \sum_{s', r} \Pr(s', r) [r + \gamma \max_{a' \in \mathcal{A}(s')} q_{\pi^*}(s', a')]$$

Select the optimal  
action locally.

follow the same optimal  
policy in the future.



# Relation to dynamic programming

Find the shortest path from the start node to the goal on a graph.


A dynamic programming algorithm will take advantage of

- optimal substructure
- overlapping subproblems

$$v_{\pi^*}(s) = \max_{a \in \mathcal{A}(s)} \sum_{s',r} \Pr(s', r | s, a) [r + \gamma v_{\pi^*}(s')]$$

Immediate “stochastic” cost:

$$-\sum_{s'} \sum_r \Pr(s', r | s, a) r$$



$$v_{\pi^*}(s')$$

**Optimal substructure**  
This must be the shortest distance from node A to the goal. Otherwise, one could have found a shorter path.

**Overlapping subproblem**  
This shortest path can be used in solving multiple larger problems.

# Machine Learning

## AIAA 5046, Spring 2024

Sihong Xie

AI Thrust

HKUST-GZ

Goals:

- Dynamic programming
- Policy iteration
- Value iteration

Readings:

RL Chap 3-6

# Dynamic programming

If the model  $\Pr(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$  is known, use DP to

- evaluate a policy
- optimize a policy


Dynamic programming: optimization (*programming*) for sequential decision making (*dynamic*).

This is not reinforcement learning yet:

- RL learns from experiences of interactions.
- RL needs to explore the environment.

Assumptions:

1. known environment dynamics.
2. optimal substructure
  - principle of optimality
3. overlapping sub-problems.



# Policy evaluation

Problem: find state-value function  $v_\pi(s)$  of a policy  $\pi$ .


Solution: iteratively compute  $v_\pi(s)$  using the Bellman equation.

```
for k=1, 2, ...
  for all state s
```

$$v_{k+1}(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} \Pr(s',r|s,a)[r + \gamma v_k(s')]$$

Synchronous update:

- maintain two arrays for the value function.
- compute the new values only after all old ones are done (two graphs for shortest path problem).
- one sweep can take long with many states.




Asynchronous update:

- update in-place: save space and more practical.
- use the latest value *immediately* to compute new values (one graph for shortest path problem).
- it is the foundation of Monte Carlo methods.

# Policy evaluation

Example




$R_t = -1$   
on all transitions

$k=0$

$v_k$  for the  
random policy

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0


greedy policy  
w.r.t.  $v_k$



random  
policy


$k=1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0



$k=2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0



Executing the following for each state per iteration:

$$v_{k+1}(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} \Pr(s', r | s, a) [r + \gamma v_k(s')]$$

↑  
Assumed deterministic

# Policy evaluation

Example (continued)

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0




$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

Converged value function  
for the random policy:  
each value is the negative of  
the expected number of steps  
to the absorbing states.

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



An optimal policy  
emerges earlier on.

optimal  
policy

# Policy improvement

After evaluating a policy  $\pi$ , the next step is to upgrade it to a better policy  $\pi'$ .

Assume deterministic policies  $a = \pi(s)$ .

Policy improvement will find the action

$$a \triangleq \pi'(s) = \arg \max_a q_\pi(s, a) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]$$

$$\implies q_\pi(s, \pi'(s)) = \max_a q_\pi(s, a) \geq \sum_a \pi(a | s) q_\pi(s, a) = v_\pi(s).$$


Environment dynamics must be known.

This is the value function  
of the previous policy  $\pi$ .

This update can be done at all state simultaneously.

# Policy improvement

It can be proved that the new policy  $\pi'$  has a better state-value function than  $\pi$ .



$$\begin{aligned} v_\pi(s) &\leq q_\pi(s, \pi'(s)) \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = \pi'(s)] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma q_\pi(S_{t+1}, \pi'(S_{t+1})) | S_t = s] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma \mathbb{E}_{\pi'}[R_{t+2} + \gamma v_\pi(S_{t+2}) | S_{t+1}, A_{t+1} = \pi'(S_{t+1})] | S_t = s] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 v_\pi(S_{t+2}) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 q_\pi(S_{t+2}, \pi'(S_{t+2})) | S_t = s] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v_\pi(S_{t+3}) | S_t = s] \\ &\leq \dots \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] \\ &= v_{\pi'}(s). \end{aligned}$$

# Policy iteration

Given an improved policy  $\pi'$ , we can repeat the evaluation-improvement cycle to obtain a better policy  $\pi''$ , until convergence:

$$\pi_0 \xrightarrow{E} v_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} v_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} v_*.$$

Evaluation: given a policy, evaluate its value function.

Improvement: given an updated value function, select a greedy policy.


- Policy improvement theorem guarantees better policies.
- Contraction Mapping Theorem guarantees convergence.

At convergence where  $v_\pi(s)$  does not change, we must have the Bellman optimality equation:

$$v_{\pi^*}(s) = \max_{a \in \mathcal{A}(s)} \sum_{s',r} \Pr(s',r|s,a)[r + \gamma v_{\pi^*}(s')]$$

# Generalized Policy Improvement

Generalized Policy Improvement (GPI):  
any iterative and alternative policy evaluation and policy improvement.



Almost all reinforcement learning methods can be described as GPI.

# Policy iteration

The policy evaluation step in policy iteration can take too long to converge. Policy improvement is possible even with a rough estimation of the value function.

Principle of optimality

$$v_{\pi^*}(s) = \max_{a \in \mathcal{A}(s)} \sum_{s', r} \Pr(s', r | s, a) [r + \gamma v_{\pi^*}(s')]$$

↑  
Improve  $v_{\pi^*}(s)$       ↓  
Evaluating  $q_{\pi^*}(s, a)$

Turned into synchronous policy iteration equation:

$$\begin{aligned} v_{k+1}(s) &\leftarrow \max_a \mathbb{E}[R_{t+1} + \gamma v_k(S_{t+1}) | S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} \Pr(s', r | s, a) [r + \gamma v_k(s')]. \end{aligned}$$

↑  
May have not converged yet.

Turned into asynchronous policy iteration equation

$$v(s) = \max_a \sum_{s', r} \Pr(s', r | s, a) [r + \gamma v(s')].$$

# Machine Learning

## AIAA 5046, Spring 2024

Sihong Xie

AI Thrust

HKUST-GZ

Goals:

- Monte Carlo methods

Readings:

RL Chap 3-6

# MC methods

Starting from here on, we are officially on reinforcement **learning**.

The dynamics  $\Pr(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$  is unavailable:

- blackjack: the player does not know the distribution of the cards;
- portfolio management (finance): market dynamics is unknown;
- Go games: unknown opponent's action that leads to the next state.

Monte Carlo methods estimate value functions using **data** (experiences / trajectories).

- agent interacts with the environment following some policy.
- the environment responds with a reward and a next state – a sample of the dynamics.


# MC methods

Monte Carlo is a computational tools to compute quantities that are hard to calculate in closed form.

Estimate the constant  $\pi$ :

1. draw a circle with radius = 1
2. generate  $n$  points in the square (area = 4) uniformly.
3. ratio of areas =  $\pi/4 \approx N_{\text{inside}} / N_{\text{outside}}$

The larger the  $n$ , the more accurate the estimation.



# MC methods

Want to estimate the expectation  $\mathbb{E}[X]$  using Monte Carlo:  $\mathbb{E}[X] \approx \frac{1}{m} \sum_{i=1}^m x^{(i)}$

Why?

- The distribution of  $X$  is unknown or hard to compute; or
- The expectation involves integrations of continuous variable.

For reinforcement learning, without the dynamics  $\Pr(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$   
the state and action value functions are hard to compute.


- $v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s]$  We need them for policy improvement.
- $q_\pi(s, a) = \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s, A_t = a].$

# MC prediction (policy evaluation)

**First-visit MC prediction, for estimating  $V \approx v_\pi$**

Input: a policy  $\pi$  to be evaluated  
 Initialize:  
 $V(s) \in \mathbb{R}$ , arbitrarily, for all  $s \in \mathcal{S}$   
 $Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$

Loop forever (for each episode):  
 Generate an episode following  $\pi$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$  ← Sampled under  $\pi$  and the unknown  $\Pr(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$   
 $G \leftarrow 0$   
 Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :  
 $G \leftarrow \gamma G + R_{t+1}$   
 Unless  $S_t$  appears in  $S_0, S_1, \dots, S_{t-1}$ : ← Count first-visits only.  
 Append  $G$  to  $Returns(S_t)$   
 $V(S_t) \leftarrow \text{average}(Returns(S_t))$  ← Approximate the expectation, using average.



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
 on all transitions

Sample trajectories:  
 2, Left, -1, 1, Left, -1  
 6, Left, -1, 5, Up, -1, 1, Left, -1  
 14, Up, -1, 10, Down, -1, 14, Right, -1

# MC prediction (policy evaluation)

First-visit MC prediction, for estimating  $V \approx v_\pi$

Input: a policy  $\pi$  to be evaluated

Initialize:

$V(s) \in \mathbb{R}$ , arbitrarily, for all  $s \in \mathcal{S}$

$Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$

Loop forever (for each episode):

Generate an episode following  $\pi$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :


$G \leftarrow \gamma G + R_{t+1}$

~~Unless  $S_t$  appears in  $S_0, S_1, \dots, S_{t-1}$ :~~

count every visit.

Append  $G$  to  $Returns(S_t)$

$V(S_t) \leftarrow$  average( $Returns(S_t)$ )



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
on all transitions

Sample trajectories:

2, Left, -1, 1, Left, -1


6, Left, -1, 5, Up, -1, 1, Up, -1

14, Up, -1, 10, Down, -1, 14, Right, -1

# MC prediction (policy evaluation)

## Remarks

- advantage:
  - no exact environment dynamics is needed;
- disadvantages:
  - works only on episodic MDPs (finite  $T$ ).
  - needs to wait until an episode ends to compute  $G$ .  
(in temporal difference, we relax these disadvantages).



# MC control (policy optimization)

Use the Generalized Policy Iteration (GPI)


$$\pi_0 \xrightarrow{E} q_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} q_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} q_*.$$

MC prediction for the E-step:

- don't need to wait for evaluation to converge: too many trajectories are needed.
- have to work with the action-value function  $q_\pi(s, a)$ , since one needs to ...

In the I-step, pick the optimal action at each state greedily to improve the policy.


- $a = \max_{a'} q_\pi(s, a')$
- Problematic: some action may need more data to get its true value reliably, and too many actions in reality



# Exploration

Without visiting a state or a state-action pair, its value can't be estimated.

- Need **exploration** to visit all states or state-action pairs to evaluate value functions.



	1	2	3
4	5	6	7
8		10	11
12	13	14	

$R_t = -1$   
on all transitions

Sample trajectories:

2, Left, -1, 1, Left, -1, 0

6, Left, -1, 5, Up, -1, 1, Up, -1, 0

14, Up, -1, 10, Down, -1, 14, Right, -1, 0

Cannot estimate  $q_\pi(1, \text{right})$  since it is not sampled.

- How about starting from each of all states? Not easy for large state spaces.
- More importantly, some states (e.g., walls for a robot) are dangerous to visit.

# MC control with exploring starts (ES)

Monte Carlo ES (Exploring Starts), for estimating  $\pi \approx \pi_*$

Initialize:

$\pi(s) \in \mathcal{A}(s)$  (arbitrarily), for all  $s \in \mathcal{S}$

$Q(s, a) \in \mathbb{R}$  (arbitrarily), for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Loop forever (for each episode):

Choose  $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$  randomly such that all pairs have probability  $> 0$

Generate an episode from  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :


$G \leftarrow \gamma G + R_{t+1}$

Unless the pair  $S_t, A_t$  appears in  $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$ :

Append  $G$  to  $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$\pi(S_t) \leftarrow \arg \max_a Q(S_t, a)$



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
on all transitions

Sample trajectories:

2, Left, -1, 1, Left, -1, 0

6, Left, -1, 5, Up, -1, 1, Up, -1, 0

14, Up, -1, 10, Down, -1, 14, Right, -1, 0

ES is not realistic for large state-action space, since it needs to start from every (s, a) pair.

# MC control with $\epsilon$ -greedy exploration

$\epsilon$ -greedy policies:  $\pi'(a|s) = \begin{cases} \frac{\epsilon}{|\mathcal{A}|} + (1 - \epsilon) & \text{if } a = \arg \max_{a'} q_\pi(a'|s), \\ \frac{\epsilon}{|\mathcal{A}|} & \text{for other actions.} \end{cases}$



	1	2	3
4	5	6	7
8		10	11
12	13	14	

$R_t = -1$   
on all transitions

So long as all states are visited, there is a non-zero probability that all  $(s, a)$  pairs will be visited in the sampled trajectories.

# MC control with $\epsilon$ -greedy exploration

An  $\epsilon$ -greedy policy improves the prior policy.

For any state  $s$ , the following inequality holds:

$$\begin{aligned}\mathbb{E}_{\pi'}[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s] &= \sum_a \pi'(a|s) q_\pi(a, s) \\ &= \frac{\epsilon}{|\mathcal{A}|} \sum_a q_\pi(s, a) + (1 - \epsilon) \max_a q_\pi(s, a) \\ &\geq \frac{\epsilon}{|\mathcal{A}|} \sum_a q_\pi(s, a) + (1 - \epsilon) \sum_a \frac{\pi(a|s) - \epsilon/|\mathcal{A}|}{1 - \epsilon} q_\pi(s, a) \\ &= \sum_a \pi(a|s) q_\pi(s, a) \\ &= v_\pi(s).\end{aligned}$$

# MC control with $\varepsilon$ -greedy exploration

On-policy first-visit MC control (for  $\varepsilon$ -soft policies), estimates  $\pi \approx \pi_*$

Algorithm parameter: small  $\varepsilon > 0$

Initialize:

$\pi \leftarrow$  an arbitrary  $\varepsilon$ -soft policy

$Q(s, a) \in \mathbb{R}$  (arbitrarily), for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following  $\pi$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair  $S_t, A_t$  appears in  $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$ :

Append  $G$  to  $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \arg \max_a Q(S_t, a)$  (with ties broken arbitrarily)

For all  $a \in \mathcal{A}(S_t)$ :

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

$\varepsilon$ -greedy

Remarks:

- It is an on-policy MC control algorithm.
  - The policies being optimized and used to select actions are the same.
- $\varepsilon$  must converge to 0 to obtain an optimal policy.
  - But if  $\varepsilon$  goes to 0 too fast, one loses exploration capability.

# Off-policy MC control



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
on all transitions

On-policy control can lack exploration capability

- $\epsilon$  may go to 0 too soon before (almost) all state-action pairs are visited.

Use a **behavior policy**  $b(a|s)$  that keeps on exploring

while learning the **target policy**  $\pi(a|s)$

This is called “off-policy” reinforcement learning: can use data collected before.

Sample trajectories from  $b(a|s)$  are biased samples of the distribution from  $\pi(a|s)$

$$\sum_{i=1}^m G_t^{(i)} \approx \mathbb{E}_b[G_t | S_t = s, A_t = a] \neq \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$

# Importance sampling (IS)

Estimate the mean of a random variable under one distribution,  
using sample from another distribution.

This holds true always:

$$\sum_x P(x)x = \mathbb{E}_P[X] = \mathbb{E}_Q \left[ \frac{P(X)}{Q(X)} X \right] = \sum_x Q(x) \frac{P(x)}{Q(x)} x$$

Monte Carlo version:

$$\sum_{i=1}^m x^{(i)} \approx \mathbb{E}_P[X] = \mathbb{E}_Q \left[ \frac{P(X)}{Q(X)} X \right] \approx \sum_{i=1}^m \frac{P(x)}{Q(x)} x^{(i)}$$

Sample from  $P$

Sample from  $Q$

# Importance sampling (IS)

Probability of a trajectory generated by  $\pi(a|s)$

$$\Pr(S_t, A_t, R_{t+1}, S_{t+1}, \dots, S_{T-1}, A_{T-1}, R_T) = \prod_{k=t}^{T-1} \pi(A_t|S_t) \Pr(S_{t+1}, R_{t+1}|S_t, A_t)$$

Probability of a trajectory generated by  $b(a|s)$

$$\Pr(S_t, A_t, R_{t+1}, S_{t+1}, \dots, S_{T-1}, A_{T-1}, R_T) = \prod_{k=t}^{T-1} b(A_t|S_t) \Pr(S_{t+1}, R_{t+1}|S_t, A_t)$$

Weight of the trajectory  $\rho_{t:T-1} = \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}$

The product of multiple probabilities leads to high variance: IS is more common in TD (discussed later) rather than MC.

# Incremental update of value functions

Assuming action-value function is estimated using  $n$  returns;

with the  $(n + 1)$ -th return, the value function can be updated as:

$$\begin{aligned} q_{\pi}^{(n+1)}(s, a) &\approx \frac{1}{n+1} \sum_{i=1}^{n+1} g_{t(i)} \\ &= \frac{1}{n+1} \left( \sum_{i=1}^n g_{t(i)} + g_{t(n+1)} \right) \\ &= \frac{n}{n+1} q_{\pi}^{(n)}(s, a) + \frac{g_{t(n+1)}}{n+1} \\ &= \left(1 - \frac{1}{n+1}\right) q_{\pi}^{(n)}(s, a) + \frac{g_{t(n+1)}}{n+1} \\ &= q_{\pi}^{(n)}(s, a) + \frac{1}{n+1} [g_{t(n+1)} - q_{\pi}^{(n)}(s, a)] \end{aligned}$$

On the left, each observation has weight 1.

More generally, each return can be associated with a weight.

We will use the following general formula:

$$q_{\pi}(S_t, A_t) \leftarrow q_{\pi}(S_t, A_t) + \alpha[G_t - q_{\pi}(S_t, A_t)]$$

# Off-policy MC control with IS

## Off-policy MC control, for estimating $\pi \approx \pi_*$

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$$Q(s, a) \in \mathbb{R} \text{ (arbitrarily)}$$

$$C(s, a) \leftarrow 0$$

$$\pi(s) \leftarrow \arg \max_a Q(s, a) \quad (\text{with ties broken consistently})$$

Loop forever (for each episode):

$b \leftarrow$  any soft policy

Generate an episode using  $b$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$$G \leftarrow 0$$

$$W \leftarrow 1$$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$$G \leftarrow \gamma G + R_{t+1}$$

$$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$$

$$\pi(S_t) \leftarrow \arg \max_a Q(S_t, a) \quad (\text{with ties broken consistently})$$

If  $A_t \neq \pi(S_t)$  then exit inner Loop (proceed to next episode)

$$W \leftarrow W \frac{1}{b(A_t | S_t)}$$

$$q_\pi(s, a) \leftarrow$$

$$q_\pi(s, a) + \alpha([\text{target value}] - q_\pi(s, a))$$

## Weighted IS

$$q_\pi(s, a) \approx \frac{\sum_{t \in \mathcal{T}(s, a)} \rho_{t:T(t)-1} G_t}{\sum_{t \in \mathcal{T}(s, a)} \rho_{t:T(t)-1}}.$$

## Incremental update

Greedy is good enough and exploration handled by policy  $b$ .

Where is  $\pi(A_t | S_t)$ ? It is already in the sampling distribution.

# Machine Learning

## AIAA 5046, Spring 2024

Sihong Xie

AI Thrust

HKUST-GZ

Goals:

- Temporal difference
- Function approximation

Readings:

RL Chap 3-6


# Temporal difference

Use the difference in the value estimation between two consecutive steps for updating value functions.

- does not need to wait until an episode to finish.
- can work with continuous (non-episodic) MDPs.
- in practice converges faster, as updates are seen immediately

In MC, we have  $q_{\pi}(S_t, A_t) \leftarrow q_{\pi}(S_t, A_t) + \alpha[G_t - q_{\pi}(S_t, A_t)]$

In TD, we have  $q_{\pi}(S_t, A_t) \leftarrow q_{\pi}(S_t, A_t) + \alpha[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) - q_{\pi}(S_t, A_t)]$



# Temporal difference


Example: given 8 trajectories sampled from an unknown MDP:

- Each state has only one action. State B: the environment has stochastic response.

Evaluate the policy by estimating the state-value function.

B, 0	B, 1
B, 1	B, 1
B, 1	B, 1
B, 1	A, 0, B, 0

TD also estimates this MDP:



- First-visit MC prediction:  $q_{\pi}(B) = \frac{6}{8}; q_{\pi}(A) = 0.$
- TD prediction:  $q_{\pi}(B) = \frac{6}{8}; q_{\pi}(A) = \frac{6}{8}$ . (assuming learning rate and discount factor are both 1).
  - Information of action–value function propagate faster.

# TD control: SARSA

On-policy TD control: learn from transitions  $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1})$

## Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A'$ ;

    until  $S$  is terminal

need exploration to visit all  $(s, a)$  pairs.

$\varepsilon$  must converge to 0  
to find an optimal policy.

on-policy

# TD control: Q-learning

## Off-policy TD control:

- Behavior policy:  $\epsilon$ -greedy policy derived from  $Q$
  - Target policy: the greedy policy derived from  $Q$

Why there is no  
importance sampling weight?  
*Behavior policy is deterministic  
and  $\pi(a^*|s) = 1$ .*

**Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$**

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize  $S$

Loop for each step of episode:

Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)  $\leftarrow$  the behavior policy

Take action  $A$ , observe  $R, S'$


$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$

until  $S$  is terminal

the target greedy policy


# Summary



# RL with function approximation

Motivations:

- Very large state spaces
  - Go game has  $10^{170}$  states
- Tabular methods can't generalized.



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$$R_t = -1 \text{ on all transitions}$$

Tabular methods use tables to keep track of value functions


s	v(s)
1	0.1
2	5
...	...

(s, a)	q(s, a)
(1, Left)	0.1
(1, Right)	5
...	...

# RL with function approximation

Approximate the value function using some supervised learning models

- capture the input-output relationships;
- can be trained as more input-output are observed;
- should support online **learning** and **forgetting**.



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
on all transitions

what tabular methods will do:

$$v(S_t) \leftarrow R_{t+1} + \gamma v(S_{t+1})$$



s	v(s)
1	-1
2	-2
...	...

Sample trajectories:

2, Left, -1, 1, Left, -1, 0

6, Left, -1, 5, Up, -1, 1, Up, -1, 0

14, Up, -1, 10, Down, -1, 14, Right, -1, 0

Function approximation methods will minimize the MSE:

$$\sum_t [R_{t+1} + \gamma \hat{v}_\pi(S_{t+1}; \mathbf{w}) - \hat{v}_\pi(S_t; \mathbf{w})]^2$$

# Online learning and forgetting

Episodes are generated during interaction


Stochastic gradient descent with MC:

$$\begin{aligned}\mathbf{w} &\leftarrow \mathbf{w} - \alpha \nabla_{\mathbf{w}} [v_{\pi}(S_t) - \hat{v}_{\pi}(S_t; \mathbf{w})]^2 \\ &= \mathbf{w} + \alpha [G_t - \hat{v}_{\pi}(S_t; \mathbf{w})] \nabla_{\mathbf{w}} \hat{v}_{\pi}(S_t; \mathbf{w})\end{aligned}$$

$G_t$  depends on the value function parameters too.

Stochastic **semi**-gradient descent with TD:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R_{t+1} + \gamma \hat{v}_{\pi}(S_{t+1}; \mathbf{w}) - \hat{v}_{\pi}(S_t; \mathbf{w})] \nabla_{\mathbf{w}} \hat{v}_{\pi}(S_t; \mathbf{w})$$



Issues: non-stationary training data

- dependencies between two consecutive states.
- in GPI, the policy keeps changing.

Some degree of forgetting is necessary.

# Linear models

Using linear model  $\hat{v}_\pi(s; \mathbf{w}) = \mathbf{w}^\top \mathbf{x}(s)$

No need to optimize the parameter so that

Stochastic gradient descent with MC:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha[G_t - \hat{v}_\pi(S_t; \mathbf{w})]\mathbf{x}(S_t)$$

$$G_t = \hat{v}_\pi(S_t; \mathbf{w})$$

$$R_{t+1} + \gamma \hat{v}_\pi(S_{t+1}; \mathbf{w}) = \hat{v}_\pi(S_t; \mathbf{w})$$

since:

- training data are non-stationary;
- may increase errors in other states.

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha[R_{t+1} + \gamma \hat{v}_\pi(S_{t+1}; \mathbf{w}) - \hat{v}_\pi(S_t; \mathbf{w})]\mathbf{x}(S_t)$$

# Control with function approximation

Need to evaluate the action-state value function by minimizing

$$\sum_{s \in \mathcal{S}, a \in \mathcal{A}} \mu_\pi(s, a) [q_\pi(s, a) - \hat{q}_\pi(s, a; \mathbf{w})]^2$$

Stochastic gradient descent with MC:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [\mathbf{G}_t - \hat{q}_\pi(S_t, A_t; \mathbf{w})] \nabla \hat{q}(S_t, A_t; \mathbf{w})$$

Stochastic gradient descent with on-policy TD (SARSA):


$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [\mathbf{R}_{t+1} + \gamma \hat{q}_\pi(S_{t+1}, A_{t+1}; \mathbf{w}) - \hat{q}_\pi(S_t, A_t; \mathbf{w})] \nabla \hat{q}_\pi(S_t, A_t; \mathbf{w})$$

Stochastic gradient descent with off-policy TD (Q-learning):

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [\mathbf{R}_{t+1} + \gamma \max_a \hat{q}_\pi(S_{t+1}, a; \mathbf{w}) - \hat{q}_\pi(S_t, A_t; \mathbf{w})] \nabla \hat{q}_\pi(S_t, A_t; \mathbf{w})$$

# Deep Q-network (DQN)

Fitting the action-value function using a neural network is not new.



$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R_{t+1} + \gamma \max_a \hat{q}_\pi(S_{t+1}, a; \mathbf{w}) - \hat{q}_\pi(S_t, A_t; \mathbf{w})] \nabla \hat{q}_\pi(S_t, A_t; \mathbf{w})$$

# Deep Q-network (DQN)

Two challenges

- data dependencies => sample mini-batch from a large buffer of transitions.
- non-stationary target

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R_{t+1} + \gamma \max_a \hat{q}_\pi(S_{t+1}, a; \mathbf{w}) - \hat{q}_\pi(S_t, A_t; \mathbf{w})] \nabla \hat{q}_\pi(S_t, A_t; \mathbf{w})$$

w keeps moving

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R_{t+1} + \gamma \max_a \hat{q}_\pi(S_{t+1}, a; \mathbf{w}^{-1}) - \hat{q}_\pi(S_t, A_t; \mathbf{w})] \nabla_{\mathbf{w}} \hat{q}_\pi(S_t, A_t; \mathbf{w})$$

Target Q-network that is  
not updated too frequent.

# Deep Q-network (DQN)

## **Algorithm 1: deep Q-learning with experience replay.**

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

**For** episode = 1,  $M$  **do**

    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

**For**  $t = 1, T$  **do**

        With probability  $\varepsilon$  select a random action  $a_t$

        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$

        Every  $C$  steps reset  $\hat{Q} = Q$

**End For**

**End For**