



# Distributional Reinforcement Learning



# Why Distributional Reinforcement Learning?

## A toy example

cliff walking with noise – wind could blow the agent to a random direction



**Left:** The Cliffs environment, along with the path preferred by the quick and safe policies. **Right:** The return distribution for these policies, estimated by sampling 100,000 trajectories from the environment. The figure reports the distribution as a probability density function computed using kernel density estimation (with bandwidth 0.02).

**Conclusion:** A quick policy that walks along the cliff's edge and a safe policy that walks two cells away from the edge. The return distribution of the faster policy, in particular, is sharply peaked around  $-1$  and  $1$ : the goal may be reached quickly, but the agent is more likely to fall.

**Advantage:** The expected return alone fails to differentiate these policies. For risk-sensitive decision-making, the safe policy is superior as its distribution avoids the significant negative tail.

# Standard & Distributional RL and Convergence

Standard RL algorithms estimate the expected value of return  $Z^\pi$ :

$$V^\pi(x) := \mathbb{E}[Z^\pi(x)] = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t R(x_t, a_t) \mid x_0 = x\right]$$

the action is fixed to a particular choice

$$Q^\pi(x, a) := \mathbb{E}[Z^\pi(x, a)] = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t R(x_t, a_t)\right]$$

$$x_t \sim P(\cdot|x_{t-1}, a_{t-1}), a_t \sim \pi(\cdot|x_t), x_0 = x, a_0 = a.$$

Standard Bellman Operator:

$$\mathcal{T}^\pi Q(x, a) = \underbrace{\mathbb{E}[R(x, a)]}_{\text{Expectation of current reward}} + \gamma \underbrace{\mathbb{E}_{P, \pi}[Q(x', a')]}_{\substack{\text{Future-to-current discount} \\ \text{expectation of future rewards}}}.$$

Distributional Bellman Operator:

$$\mathcal{T}^\pi Z(x, a) \stackrel{D}{=} \underbrace{R(x, a)}_{\text{Distribution of current reward}} + \gamma \underbrace{Z(x', a')}_{\text{Distribution of future rewards}},$$

$$x' \sim P(\cdot|x, a), a' \sim \pi(\cdot|x')$$

Bellman optimality operator:

$$Q(x, a) = \mathcal{T}Q(x, a) := \mathbb{E}[R(x, a)] + \gamma \mathbb{E}_P \max_{a'} Q(x', a')$$

Distributional Bellman optimality operator:

$$\mathcal{T}Z(x, a) \stackrel{D}{=} R(x, a) + \gamma Z(x', \arg \max_{a' \in \mathcal{A}} \mathbb{E} Z(x', a'))$$

Standard TD update:

$$V(x) \leftarrow V(x) + \alpha(r + \gamma V(x') - V(x)),$$

$$a \sim \pi(\cdot|x), r \sim R(x, a), x' \sim P(\cdot|x, a).$$

What is distributional TD update?


# Return Distribution Representation

**Q:** How to represent the return distribution and how to do the dynamic programming?

**A:** Network architectures for DQN and recent distributional RL algorithms.

**Return Representation: given a distribution by taking a fixed action**  $Z(s, a) \sim \mathcal{N}(\mu = 10, \sigma^2 = 4)$

## Representation Way



## Expectation

Predefined  
uniform  
return atoms

$$z = [0, 5, 10, 15, 20] \quad \tau = [0.1, 0.3, 0.5, 0.7, 0.9]$$

$$Z_{C51}(s, a) = \{(z_i, p_i)\}_{i=1}^N$$

Predefined  
uniform return  
quantiles


$$\theta_i(s, a) \approx F_Z^{-1}(\tau_i)$$

## Sampled quantiles


$$\tau = 0.13, 0.47, 0.88, \dots$$

$$f_\theta(s, a, \tau) \approx F_Z^{-1}(\tau)$$


## DQN



## C51



## QR-DQN





## IQN


# Categorical and Quantile-based Representation

**Q: How to represent the return distribution and how to do the dynamic programming?**

**A1: Categorical (Predefined uniform return atoms):**



- (a)  $P^\pi Z = \{(0, 0.01), (5, 0.20), (10, 0.58), (15, 0.20), (20, 0.01)\}$
- (b)  $\gamma P^\pi Z \rightarrow \{(0, 0.01), (4.5, 0.20), (9, 0.58), (13.5, 0.20), (18, 0.01)\}$
- (c)  $R + \gamma P^\pi Z \rightarrow \{(2, 0.01), (6.5, 0.20), (11, 0.58), (15.5, 0.20), (20, 0.01)\}$
- (d)  $\Phi \mathcal{T}^\pi Z \rightarrow \{(0, 0.006), (5, 0.144), (10, 0.524), (15, 0.296), (20, 0.030)\}$




*Figure 1.* A distributional Bellman operator with a deterministic reward function: (a) Next state distribution under policy  $\pi$ , (b) Discounting shrinks the distribution towards 0, (c) The reward shifts it, and (d) Projection step (Section 4).

# Categorical and Quantile-based Representation

**Q: How to represent the return distribution and how to do the dynamic programming?**

**A2: Quantile-based methods:**

Predefined uniform  
return quantiles



QR-DQN

Sampled quantiles



IQN

**QRDQN (with predefined quantile positions) :**

$$\begin{aligned} Q(x', a') &:= \sum_j q_j \theta_j(x', a') \\ a^* &\leftarrow \arg \max_{a'} Q(x, a') \\ \mathcal{T} \theta_j &\leftarrow r + \gamma \theta_j(x', a^*), \quad \forall j \end{aligned}$$

$\theta_i(x, a)$  represents the estimated value of the return distribution  $Z_\theta(x, a)$  at the  $i$ -th quantile.

**Objective function with quantile Huber loss**

$$\sum_{i=1}^N \mathbb{E}_j [\rho_{\hat{\tau}_i}^\kappa (\mathcal{T} \theta_j - \theta_i(x, a))]$$

$$\rho_\tau^\kappa(u) = |\tau - \delta_{\{u<0\}}| \mathcal{L}_\kappa(u). \quad \mathcal{L}_\kappa(u) = \begin{cases} \frac{1}{2}u^2, & \text{if } |u| \leq \kappa \\ \kappa(|u| - \frac{1}{2}\kappa), & \text{otherwise} \end{cases}.$$

If we underestimate the target (target > predicted;  $u>0$ ), we multiply by the weight  $\tau$ . If we overestimate, we multiply by  $1 - \tau$ . It means that if the quantile  $\tau$  is big, we penalize underestimation more to balance the data distribution over quantile.

## Robust Quadrupedal Locomotion via Risk-Averse Policy Learning


Jiyuan Shi<sup>1,2\*</sup>, Chenjia Bai<sup>2†</sup>, Haoran He<sup>2,3</sup>, Lei Han<sup>4</sup>, Dong Wang<sup>2</sup>, Bin Zhao<sup>2,5</sup>,  
Mingguo Zhao<sup>1</sup>, Xiu Li<sup>1</sup>, Xuelong Li<sup>2,5</sup>

### Aleatoric Uncertainty:

The inherent, irreducible randomness in  
the environment

### Robust locomotion controller:

Encountering risks in the environment,  
such as sudden pushes and missing a step.



# Problem Definition

**MDP:**

$$\mathbf{o}_t = [\mathbf{v}_t, \boldsymbol{\omega}_t, \mathbf{g}_t, \mathbf{c}_t, \mathbf{q}_t, \dot{\mathbf{q}}_t, \mathbf{a}_{t-1}] \in \mathbb{R}^{48}$$

Action: target joint positions

**Reward:**

$$\mathbf{r} = \mathbf{r}_{\text{task}} - \sum_{i=1}^M w_i \mathbb{I}_{|s_i| > \bar{s}_i} \cdot \mathcal{B}_p$$


TABLE I: Definition of reward functions. Here  $\tau$  refers to joint torque.

Reward Term	Definition	Weight
linear velocity tracking	$e^{-(\mathbf{v}_{xy}^{\text{cmd}} - \mathbf{v}_{xy})^2 / \sigma}$	5
angular velocity tracking	$e^{-(\boldsymbol{\omega}_{\text{yaw}}^{\text{cmd}} - \boldsymbol{\omega}_{\text{yaw}})^2 / \sigma}$	0.5
linear velocity penalty	$v_z^2$	-1.0
angular velocity penalty	$\boldsymbol{\omega}_{\text{roll,pitch}}^2$	-0.05
joint acceleration	$\dot{\mathbf{q}}^2$	-2.5e-7
torques	$\tau^2$	-2e-5
action magnitude	$\mathbf{a}^2$	-0.01
collision	$n_{\text{collision}}$	-1e-3
action rate	$(\mathbf{a}_t - \mathbf{a}_{t-1})^2$	-0.01
torque smooth	$(\tau_t - \tau_{t-1})^2$	-3e-4
feet air time	$\sum_{f=0}^4 (\mathbf{t}_{\text{air},f} - 0.5)$	2

TABLE II: Definition of risks. The risk terms will be added to the reward function defined in (2).

Risk Term	Threshold	Weight
base pitch	0.5 rad	20
base roll	1 rad	100
joint velocity	10 rad·s <sup>-1</sup>	100
joint acceleration	1000 rad · s <sup>-2</sup>	100
joint torque	40 N·m	150

# Method Framework



**Risk-sensitive policy learning:** Critic is based on IQN. The goal of RALL is to find a risk-sensitive policy by maximizing the distorted expectation of  $Z$ . (V没用到, beta, At; intuiation)

$$\nabla_\phi J(\phi) = \mathbb{E} \left[ \sum_{t=0}^T A_t \nabla_\phi \log (\pi_\phi(a_t | s_t)) \right]$$

$$V_\beta(s) := \mathbb{E}_{\tau \sim U([0,1])} [Z_{\beta(\tau)}(s)]$$

$$\text{CVaR}_\eta(Z(s)) = \mathbb{E} [Z_{\tau \sim U([0,\eta])}(s)]$$

$$\nabla_\phi J(\phi) = \mathbb{E} \left[ \sum_{t=0}^T A_t^{\text{CVaR}} \nabla_\phi \log(\pi_\phi(a_t | s_t)) \right]$$

$$A_t^{\text{CVaR}} = r_t + \gamma \text{CVaR}_\eta(Z(s_{t+1})) - \text{CVaR}_\eta(Z(s_t))$$

**Aleatoric uncertainty of return:**

$$IQR = Q_3 - Q_1, \quad Q_3 = F_Z^{-1}(0.75), \quad Q_1 = F_Z^{-1}(0.25)$$

When IQR is higher than a threshold, choose the policy trained by a lower CVaR value (0.5)