



# Introduction to NLP

AIAA 4051 Lecture 24

Instructor: Sihong Xie

Email: [sihongxie@hkust-gz.edu.cn](mailto:sihongxie@hkust-gz.edu.cn)

AI Thrust, HKUST(GZ)

# Score-Based Perspective: The Score Function

An alternative view of Diffusion comes from Score-based modeling. Instead of probability densities, we learn the Score Function:  $s(x) = \nabla_x \log p(x)$

This vector field points toward the high-density regions of the data distribution (the "peaks"). Denoising is essentially the act of moving a noisy sample back toward the high-probability manifold.

## Score Function

$$s(x) = \nabla_x \log p(x) \\ \in \mathbb{R}^d$$


## Interpretation

$s(x)$  points in the direction of steepest increase of log-density.

Arrows point toward the high-density regions (modes).

## Denoising as Gradient Ascent


Starting from noise, we iteratively move samples in the direction of  $s(x)$  to reach the data manifold.



# Score-Based Perspective: The Score Function

Langevin Dynamics describes the trajectory of a particle  $x$  moving in a potential field  $U(x) = -\log p(x)$  towards the point with the lowest potential, with stochastic thermal fluctuations.

Equivalent to maximize the  $p(x)$ , which can be the real data distribution, thus generating realistic data.



In practice, we do not know  $p(x)$ .

Train a neural network to approximate the score function:

$$s_{\theta}(x) \approx \nabla_x \log p(x) \quad \Rightarrow \quad \mathcal{L}_{score} = \mathbb{E}_{p(x)} [\|s_{\theta}(x) - \nabla_x \log p(x)\|^2]$$

# Score-Based Perspective: The Score Function

To solve the “unknown score” issue, we perturb the data  $x_0$  with noise to create  $x_t$ . The log-gradient of the conditional distribution is easy to calculate:

$$\log q(x_t|x_0) = -\frac{1}{2\sigma^2}\|x_t - x_0\|^2 + \text{const} \quad \longrightarrow \quad \nabla_{x_t} \log q(x_t|x_0) = -\frac{x_t - x_0}{\sigma^2}$$

The gradient of the conditional density provides an unbiased estimate of the gradient of the marginal density.

Denoising Score Matching:  $\mathbb{E}_{q(x_0, x_t)}[\|s_\theta(x_t) - \nabla_{x_t} \log q(x_t|x_0)\|^2] = \mathbb{E}_{p(x_t)}[\|s_\theta(x_t) - \nabla_{x_t} \log p(x_t)\|^2]$

1. We know  $\nabla_{x_t} \log q(x_t|x_0) = -\frac{x_t - x_0}{\sigma^2}$ .
2. From the noise equation  $x_t = x_0 + \sigma\epsilon$ , we have  $(x_t - x_0) = \sigma\epsilon$ .
3. Substituting this:  $\nabla_{x_t} \log q(x_t|x_0) = -\frac{\sigma\epsilon}{\sigma^2} = -\frac{\epsilon}{\sigma}$ .



$$\mathcal{L}_{DSM} = \mathbb{E}_{x_0, \epsilon} \left[ \left\| s_\theta(x_t) - \left( -\frac{\epsilon}{\sigma} \right) \right\|^2 \right]$$

# Deriving

In score-based modeling and diffusion processes,  $x_t = x_0 + \sigma\epsilon$ ,  $\epsilon \sim \mathcal{N}(0, I)$

This implies that the conditional distribution:

$$q(x_t|x_0) = \frac{1}{\sqrt{(2\pi\sigma^2)^d}} \exp\left(-\frac{\|x_t - x_0\|^2}{2\sigma^2}\right)$$

So:

$$\log q(x_t|x_0) = -\frac{1}{2\sigma^2}\|x_t - x_0\|^2 + \text{const}$$

The first term is a constant with respect to  $x_t$

$$\log q(x_t|x_0) = \log\left(\frac{1}{\sqrt{(2\pi\sigma^2)^d}}\right) - \frac{\|x_t - x_0\|^2}{2\sigma^2}$$

Obtaining the score function:


$$\nabla_{x_t} \log q(x_t|x_0) = -\frac{x_t - x_0}{\sigma^2}$$

Unbiasedness, refer to [Vincent, P.]

Vincent, P. (2011). A Connection Between Score Matching and Denoising Autoencoders. Neural Computation.

# Langevin Dynamics for inference (reverse process)

Starting from a random location  $x_T$ , use the **trained** score function  $s_\theta(x_t)$  to iteratively move  $x_T$  to  $\hat{x}_0$





# Unified Framework: Denoising vs. Score Matching

$$\epsilon_{\theta}(x_t, t) \propto -\sigma_t \nabla_x \log p_t(x)$$

| Aspect                                                                                                                                               | Diffusion Models (DDPM View)                                                                                                                                                                                                                                                                                                          | Score Matching (Score-Based View)                                                                                                              |
|------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
|  1. Sampling / Generation                                           | Reverse diffusion (DDPM):<br>$x_{t-1} = \mu_{\theta}(x_t, t) + \sigma_t z, \quad z \sim \mathcal{N}(0, I)$ <div>             Typical choice: <math display="block">\mu_{\theta}(x_t, t) = \frac{1}{\sqrt{\bar{\alpha}_t}} \left( x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(x_t, t) \right)</math> </div> | Langevin Dynamics (Euler-Maruyama):<br>$x_{t-1} = x_t + \frac{\eta_t}{2} s_{\theta}(x_t, t) + \sqrt{\eta_t} z, \quad z \sim \mathcal{N}(0, I)$ |
|  2. Dynamics Type                                                   | Discrete-time Markov chain<br>(finite $T$ steps)                                                                                                                                                                                                                                                                                      | Continuous-time SDE (infinity small step limit)                                                                                                |
|  3. Key Connection                                                | $\epsilon_{\theta}(x_t, t) = -\sigma_t \nabla_x \log p_t(x_t) + (\text{constant factor})$                                                                                                                                                                                                                                             | $\epsilon_{\theta}(x_t, t) = -\sigma_t \nabla_x \log p_t(x_t) + (\text{constant factor})$                                                      |
|  4. Takeaway                                                      | Learn to <b>denoise</b> noisy data and recover clean samples.                                                                                                                                                                                                                                                                         | Estimate the <b>score function</b> to sample from the <b>data distribution</b> .                                                               |
| ★ <b>Big Idea:</b> Both methods operate on the <b>same sequence</b> of noisy data distributions $p_t(x)$ and are equivalent up to a constant factor. |                                                                                                                                                                                                                                                                                                                                       |                                                                                                                                                |

# Why Diffusion for NLP?

## Autoregressive (GPT-style) Generation



### Task (Global Constraint):


Write a 1-line poem where the last word is "Sun".



### Challenge:

The model only sees the past.  
It cannot easily plan for the future.

### Left-to-Right Generation (Step-by-Step)



By the time it gets to the end, the context is already fixed.  
It's hard (or impossible) to end with "Sun" without awkwardness.

### The Problem: Myopic Decisions

At each step,  
GPT chooses the  
best local word...



...but gets stuck in a  
corner of the space  
that doesn't allow  
the constraint.



### Result:

Autoregressive models struggle with global constraints,  
multi-point attributes, and long-range coherence.

## Diffusion Model Generation (Whole-Sequence Refinement)



### Same Task (Global Constraint):


Write a 1-line poem where the last word is "Sun".



### Advantage:


The model sees and refines  
the entire sequence at once.

### Simultaneous Refinement (Many Steps)



### Why It Works: Global View Enables Planning

The model optimizes the  
entire sequence at every step.



It can "plan" for the future,  
ensuring the last word is "Sun"  
while keeping the sentence natural.

Global alignment ✓



### Result:

Diffusion models naturally excel at global constraints,  
multi-point control, style blending, and long-range coherence.


### Examples of Control

- ✓ Last word / first word
- ✓ Include specific words or phrases
- ✓ Follow a rhyme or meter scheme
- ✓ Match a topic and sentiment
- ✓ Avoid certain words
- ✓ Enforce entity relationships
- ✓ And many more!



# Advanced Diffusion: From Images to Text


Treating with image pixels as continuous variables are easy: the distance metric and ordering are well-defined.



sentence =  $[w_1, w_2, w_3]$ ,  $w_i \in V, i = 1, 2, 3$


There is no order or distance metric defined on  $V$

But languages, on the other hand, are not too much different from images.



These are discrete values too.

# Diffusion of text samples



# Need to be careful about text perturbations

## Key Idea

In NLP, tokens are discrete indices. Adding Gaussian noise to a word embedding moves the point into “void” space, where it may not correspond to any valid word.

### 1 Start: A valid word

The word “cat” has embedding  $e(\text{“cat”})$  in the embedding space.



### 2 Add Gaussian noise

We add noise  $\sim \mathcal{N}(0, \sigma^2 \mathbf{I})$  to the embedding.




### 3 Lands in “void” space

The noisy point is likely far from any valid token (no meaning).



## Embedding Space Visualization

- Valid word (token)
- $e(\text{“cat”})$  (original)
- Add Gaussian noise
- Noisy point (in void)
- Void space (no valid tokens)



### Intuition

In continuous data (images), every point in space is a valid (though noisy) image.

In discrete data (text), most points in space do not correspond to any valid token.

The noisy point is far from any valid token centroid. It does not represent any meaningful word.

# From Gaussian Kernels to Transition Matrices

To define diffusion over a discrete vocabulary  $V$ , we replace the Gaussian kernel with a Transition Matrix  $Q_t \in \mathbb{R}^{V \times V}$ .

Instead of shifting a mean, we define the probability of a token  $x_{t-1}$  "jumping" to another token  $x_t$ . The forward process is now a categorical Markov chain:

$x_t = Q_t x_{t-1}$ : transiting from a probability distribution  $x_{t-1}$  over tokens to another distribution  $x_t$ .

Examples:

# How to Forward Step Works

Given current token  $x_{t-1}$  with one-hot vector  $\mathbf{x}_{t-1}$

Three kinds of transitions:

## Keep:

The token remains unchanged at its current position.

In the early stages of the forward process (close to  $t=0$ ), most tokens are "kept" to maintain the original data structure.

Represents the probability  $1 - \beta_t$  that a token survives the transition step without being corrupted.

$$Q_t = (1 - \beta_t)I + \beta_t U$$

$I$  : identity matrix (keep the same token)

$U$  : uniform matrix (all entries =  $1/V$ )

$\beta_t \in [0, 1]$  : noise level at step  $t$

## Mask:

The token is replaced by a special [MASK] token, which is unknown to the model.

This is the most popular form of discrete diffusion, often referred to as "absorbing" because the [MASK] state acts as a sink.

$$Q_t = (1 - \beta_t)I + \beta_t \mathbf{e}_{mask} \mathbf{1}^T$$

•  $I$  : identity matrix (keep the same token)

•  $\mathbf{e}_{mask}$  : one-hot vector for [MASK] token

•  $\mathbf{1}$  : all-ones column vector

•  $\beta_t \in [0, 1]$  : noise level at step  $t$

## Swap:

The original token is replaced by another token chosen from the vocabulary, either uniformly or according to a pre-defined probability matrix.

This introduces noise by replacing valid data with random incorrect data, sometimes referred to as "uniform diffusion".

$$Q_t = (1 - \beta_t)I + \beta_t \frac{1}{V-1} (\mathbf{1}\mathbf{1}^T - I)$$

•  $I$  : identity matrix (keep the same token)

•  $\mathbf{1}$  : all-ones column vector


•  $V$  : vocabulary size

•  $\beta_t \in [0, 1]$  : noise level at step  $t$



# The Forward Posterior in Discrete Space

1. Forward Process in Discrete Space:



2. We define:  $\bar{Q}_t = Q_1 Q_2 \dots Q_t$ .

3. Forward Posterior via Bayes' Rule:

$$q(x_{t-1} | x_t, x_0) = \frac{q(x_t | x_{t-1}) q(x_{t-1} | x_0)}{q(x_t | x_0)}$$

4. In terms of transition matrices:

$$q(x_t | x_{t-1}) = Q_t(x_{t-1}, x_t)$$

$$q(x_{t-1} | x_0) = \bar{Q}_{t-1}(x_0, x_{t-1})$$

$$q(x_t | x_0) = \bar{Q}_t(x_0, x_t)$$

$$q(x_{t-1} | x_t, x_0) = \frac{Q_t(x_{t-1}, x_t) \bar{Q}_{t-1}(x_0, x_{t-1})}{\bar{Q}_t(x_0, x_t)}$$

5. Training the Reverse Model:

We train a model to approximate the true reverse posterior  $q(x_{t-1} | x_t, x_0)$

**Objective: Minimize KL Divergence**

$$\mathcal{L}_{\text{KL}} = \mathbb{E}_{q(x_0) q(x_t | x_0)} \left[ D_{\text{KL}} \left( q(x_{t-1} | x_t, x_0) \parallel p_\theta(x_{t-1} | x_t) \right) \right]$$

based on variational lower bound from the last lecture.

For discrete distribution, this KL is minimized by the Category Cross-Entropy loss:

**Categorical Cross-Entropy Loss**

$$\mathcal{L}_{\text{CE}} = - \mathbb{E}_{q(x_0) q(x_t | x_0)} \left[ \log p_\theta(x_{t-1} = \hat{x}_{t-1} | x_t) \right]$$

where  $\hat{x}_{t-1} \sim q(x_{t-1} | x_t, x_0)$


6. We can train a categorical predictor to reverse the corruption process.

# Diffusion in Continuous Latent Spaces

When the vocabulary is large, discrete diffusion becomes computationally expensive (matrices are  $V \times V$ ). Diffusion-LM takes a different path: map tokens to a continuous embedding space (**with order and metric!**) and apply Gaussian Diffusion.

## Token Island

States are tokens from a vocabulary of size  $V$ .




Very large  $V$   
 $\Rightarrow$  Transition  
matrix  $Q_t \in \mathbb{R}^{V \times V}$   
becomes expensive.

## The Bridge: Token Embedding

Learn an embedding function

$$E: \{1, \dots, V\} \rightarrow \mathbb{R}^d$$

that maps each token to a vector.



## Diffusion-LM (Continuous Latents)

### Vector Continent

Work in a continuous embedding space  $\mathbb{R}^d$ .



No  $V \times V$  matrix.  
Use Gaussian  
diffusion in  $\mathbb{R}^d$ .




Diffusion-LM iteratively denoises a sequence of Gaussian vectors into word vectors, yielding a intermediate latent variables of decreasing noise level  $\{\mathbf{X}_T, \dots, \mathbf{X}_0\}$

# Solution

Suppose we have a sequence of words:  $\mathbf{w} = \{w_1, w_2, \dots, w_n\}$ . An embedding function maps each word into a vector:  $Emb(w_i) \in \mathbb{R}^d$ .

Thus, the entire sequence is encoded into:  $\mathbf{x}_0 = Emb(\mathbf{w}) \in \mathbb{R}^{n \times d}$



Loss function:


$$\mathcal{L}_{simple}^{e2e}(\mathbf{w}, \theta) = \mathbb{E}_{q_\phi(\mathbf{x}_0 | \mathbf{w})} \left[ \underbrace{\mathcal{L}_{simple}(\mathbf{x}_0)}_{\text{diffusion Loss}} + \underbrace{\|Emb(\mathbf{w}) - \overbrace{\mu_\theta(\mathbf{x}_0, 1)}^{\text{predicted } \mathbf{x}_0}\|^2}_{\text{rounding}} - \log p_\theta(\mathbf{w} | \mathbf{x}_0) \right]$$

diffusion Loss:

$$\mathcal{L}_{simple}(\mathbf{x}_0, \theta) = \sum_{t=1}^T \mathbb{E}_{q(\mathbf{x}_t | \mathbf{x}_0)} [\|\mu_\theta(\mathbf{x}_t, t) - \hat{\mu}(\mathbf{x}_t, \mathbf{x}_0)\|^2],$$


] rounding to discrete tokens using softmax

$$p_\theta(\mathbf{w} | \mathbf{x}_0) = \prod_{i=1}^n p_\theta(w_i | \mathbf{x}_i)$$



# How to control our text generation?

We can now generate the correct text, but how do we generate the text we want?



We iteratively perform gradient updates on these continuous embeddings to optimize for fluency and satisfy control requirements  $\mathbf{c}$  (parametrized by a classifier).

Controlling  $\mathbf{x}_{0:T}$  is equivalent to decoding from the posterior  $p(\mathbf{x}_{0:T}|\mathbf{c}) = \prod_{t=1}^T p(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{c})$  and we decompose this joint inference problem to a sequence of control problems at each diffusion step:

$$p(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{c}) \propto p(\mathbf{x}_{t-1} | \mathbf{x}_t) \cdot p(\mathbf{c} | \mathbf{x}_{t-1}, \mathbf{x}_t) \quad (\text{Bayes Theorem})$$

$$\nabla_{\mathbf{x}_{t-1}} \log p(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{c}) = \nabla_{\mathbf{x}_{t-1}} \log p(\mathbf{x}_{t-1} | \mathbf{x}_t) + \nabla_{\mathbf{x}_{t-1}} \log p(\mathbf{c} | \mathbf{x}_{t-1}),$$

# Deriving

Controlling  $\mathbf{x}_{0:T}$  is equivalent to decoding from the posterior  $p(\mathbf{x}_{0:T}|\mathbf{c}) = \prod_{t=1}^T p(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{c})$  and we decompose this joint inference problem to a sequence of control problems at each diffusion step:

$$p(x_{t-1}|x_t, c) \propto p(x_{t-1}|x_t) \cdot p(c|x_{t-1}, x_t)$$

To simplify the second term, we adopt conditional independence assumptions from prior work on controlling diffusions:

$$p(c|x_{t-1}, x_t) \approx p(c|x_{t-1})$$

To perform gradient-based optimization, we take the logarithm of both sides:

$$\log p(x_{t-1}|x_t, c) = \log p(x_{t-1}|x_t) + \log p(c|x_{t-1}) + \text{const}$$

So:

$$\nabla_{x_{t-1}} \log p(x_{t-1}|x_t, c) = \nabla_{x_{t-1}} \log p(x_{t-1}|x_t) + \nabla_{x_{t-1}} \log p(c|x_{t-1})$$




# How to train a classifier

A labeled dataset  $\mathcal{D} = \{(w, c)\}$ , consisting of text sequences  $w$  and their corresponding control labels  $c$ .

Classifier  $\phi$

**Repeat** until convergence:

1. Sample  $(w, c) \sim \mathcal{D}$ .
2. Map text to clean embedding:  $x_0 \sim q_\phi(x_0|w) = \mathcal{N}(EMB(w), \sigma_0 I)$ .
3. Sample a random timestep:  $t \sim \text{Uniform}(\{1, \dots, T\})$ .
4. Generate noisy latent via forward process:  $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$ , where  $\epsilon \sim \mathcal{N}(0, I)$ .
5. Predict control attribute:  $\hat{p} = \text{Classifier}_\phi(x_t, t)$ .
6. Compute loss:  $\mathcal{L} = -\log \hat{p}(c)$ .
7. Update **only** classifier parameters  $\phi$  via gradient descent.




# Experiment

Example 1

Target Syntactic Parse (linearized)


( S ( S ( NP \* ) ( VP \* ( NP ( NP \*\* ) ( VP \* ( NP ( ADJP \*\* ) \* ) ) ) ) ) \* ( S ( NP \*\*\* ) ( VP \* ) ( ADJP ( ADJP \* ) ) ) ) )



Example 2

Target Syntactic Parse (linearized)

( S ( S ( VP \* ( PP \* ( NP \*\* ) ) ) ) ) \* ( NP \*\*\* ) ( VP \* ( NP ( NP \*\* ) ( SBAR ( WHNP \* ) ( S ( VP \* ( NP \*\* ) ) ) ) ) ) ) )



S: Sentence    NP: Noun Phrase    VP: Verb Phrase    PP: Prepositional Phrase  
ADJP: Adjective Phrase    SBAR: Subordinate Clause    WHNP: Wh-Noun Phrase    \*: Terminal (word)

1. This represents **two sentences (clauses) joined by a single connector word (\*)**. The first clause requires a **single-word subject (NP \*)** and a **deeply nested verb phrase**. The second clause requires a **subject that is precisely three words long (NP \*\*\*)** followed by a verb and an adjective phrase.
2. The notation includes an SBAR, which requires the model to generate a relative clause.

|                 |                                                                                                                                                                      |
|-----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Syntactic Parse | ( S ( S ( NP * ) ( VP * ( NP ( NP ** ) ( VP * ( NP ( ADJP ** ) * ) ) ) ) ) * ( S ( NP *** ) ( VP * ( ADJP ( ADJP * ) ) ) ) )                                         |
| FUDGE           | <b>Zizzi is a cheap restaurant . [incomplete]</b>                                                                                                                    |
| Diffusion-LM    | Zizzi is a pub providing <b>family friendly Indian food</b> Its customer rating is low                                                                               |
| FT              | Cocum is a Pub serving <b>moderately priced meals</b> and the customer rating is high                                                                                |
| Syntactic Parse | ( S ( S ( VP * ( PP * ( NP ** ) ) ) ) * ( NP *** ) ( VP * ( NP ( NP ** ) ( SBAR ( WHNP * ) ( S ( VP * ( NP ** ) ) ) ) ) ) ) )                                        |
| FUDGE           | <b>In the city near The Portland Arms is a coffee and fast food place named The Cricketers which is not family - friendly with a customer rating of 5 out of 5 .</b> |
| Diffusion-LM    | Located on the riverside , <b>The Rice Boat</b> is a restaurant that serves Indian food .                                                                            |
| FT              | Located near The Sorrento, <b>The Mill is a pub that serves Indian cuisine.</b>                                                                                      |

Table 3: Qualitative examples from the Syntax Tree control. The syntactic parse tree is linearized by nested brackets representing the constituents, and we use the standard PTB syntactic categories. FUDGE: Controlled text generation with future discriminators. ACL 2021 FT Fine-tuning Oracle

S: Sentence    NP: Noun Phrase    VP: Verb Phrase    PP: Prepositional Phrase    ADJP: Adjective Phrase    SBAR: Subordinate Clause    WHNP: Wh-Noun Phrase    \*: Terminal (word)

# Case Study: Block Diffusion

Discrete diffusion models are limited by fixed-length generation and the inability to reuse computation via KV-caching. Block Diffusion introduces a hybrid paradigm that interpolates between Autoregressive (AR) and Diffusion models.

## Autoregression:

✓ High quality

✓ Arbitrary-length

✓ KV caching

✗ Not Parallelizable

Generation steps

There are three categories of the average  
There are three categories of the average rate  
There are three categories of the average rate of...

## Diffusion:

✗ Lower quality

✗ Fixed-length

✗ No KV caching

✓ Parallelizable

the reusability will continue to the  
Repeal the reusability cuts and the law will continue to reduce the  
Repeal the reusability cuts and prove the law will continue to reduce the deficit.

## Block Diffusion (Ours):

✓ High quality

✓ Arbitrary-length


✓ KV caching

✓ Parallelizable

On September 17, we be  
On September 17, 2016, we will be giving the release of  
On September 17, 2016, we will be giving the beta-release of the to our server testing ...

## 1 Hierarchical Factorization

The sequence likelihood is factorized over blocks:




$$p_{\theta}(\mathbf{x}) = \prod_{b=1}^B p_{\theta}(\mathbf{x}_b \mid \mathbf{x}_{<b})$$

Each block  $\mathbf{x}_b$  is generated conditioned on all previous blocks  $\mathbf{x}_{<b}$ .

## 2 Intra-block Diffusion

For each block  $b$ , we use a discrete denoising diffusion process to sample all tokens in parallel, conditioned on the fixed KV-cache of previous blocks.



Parallel denoising within the block

Arriola, Marianne, et al. "Block diffusion: Interpolating between autoregressive and diffusion language models." ICLR 2025.



# Autoregressive vs. Diffusion

Unlike GPT, where a single "bad" word can derail the entire sentence (exposure bias), Diffusion can self-correct. In early steps, it sets the global semantics; in later steps, it fixes local syntax.

| Aspect                                                                                                 | Autoregressive (GPT-style)                                                                                                                                                     | Diffusion (Diffusion-LM)                                                                                                                                                                                               |
|--------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|  Factorization        | Causal / Left-to-Right<br>$p(x) = \prod_{t=1}^T p(x_t   x_{<t})$                                                                                                               | Non-Causal / Global<br>Model $p(x_T)$ and refine to $p(x_0)$                                                                                                                                                           |
|  Generation Speed     | Fast<br>(1 forward pass per token)                                                                                                                                             | Slow (Vanilla)<br>( $T$ forward passes, $T \gg$ sequence length)                                                                                                                                                       |
|  Information Flow     | Unidirectional<br>(Past $\rightarrow$ Future)                                                                                                                                  | Bidirectional<br>(Global context at every step)                                                                                                                                                                        |
|  Planning Capability | Weak Global Planning<br>(Hard to enforce future constraints)                                                                                                                   | Strong Global Planning<br>(Sees the whole sequence during denoising)                                                                                                                                                   |
|  Strengths          | <ul style="list-style-type: none"><li>✓ Extremely fast inference</li><li>✓ Excels at short-form, open-ended generation</li><li>✓ Huge ecosystem &amp; mature tooling</li></ul> | <ul style="list-style-type: none"><li>✓ Superior controllability (style, structure, constraints)</li><li>✓ Better coherence for long-range dependencies</li><li>✓ Self-correction &amp; iterative refinement</li></ul> |
|  Weaknesses         | <ul style="list-style-type: none"><li>✗ Exposure bias / Error accumulation</li><li>✗ Struggles with global constraints</li><li>✗ Degeneration in long sequences</li></ul>      | <ul style="list-style-type: none"><li>✗ Slow inference (though improving)</li><li>✗ Higher computational cost</li><li>✗ More complex training &amp; sampling</li></ul>                                                 |